

How Attention Impact Language Learning Tasks? A Visual Analytics Tool for Interpreting, Debugging, and Tuning Text Entailment Model

Abstract—

Index Terms—Natural Language Processing



1 INTRODUCTION

How to interpret classification task

- What feature / dimension / part-of-the-input should change to alter the outcome
- What Examples can you provide that is similar to the current instance

Related works

- Interacting with Predictions: Visual Inspection of Black-box Machine Learning Models, Use partial dependency (varying along one feature and observe how prediction changes)
- Why Should I Trust You? Explaining the Predictions of Any Classifier (LIME)

ACKNOWLEDGMENTS

The authors wish to thank A, B, and C. This work was supported in part by a grant from XYZ (# 12345-67890).